

Lect. 16: Interrupt Mechanism and Context Switching

Xin Liu

Florida State University

xliu15@fsu.edu

CIS 5370 Computer Security

<https://xinliulab.github.io/cis5370.html>

The Ideal Processor

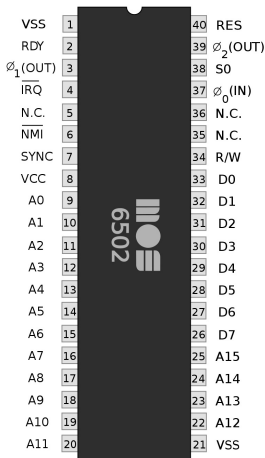
- A machine that executes instructions ****unconditionally****.

```
for (day = TODAY; day != FOREVER; day++) {  
    say("I_love_you\n");  
}
```

The Real Processor is NOT "Unconditionally Executing Instructions"

- It ****"cares"**** and responds to external ****interrupts****.
- If you fall into an infinite loop in the library...
 - A friendly security guard will "interrupt" you.

Interrupt = A Single Wire



- "Telling the processor: Stop, something has happened!"
- "The rest is up to the processor."

If the Processor Interrupts are Enabled

- **x86 Family (CISC, a legacy of history; a nightmare for processor designers)**
 - Reads interrupt vector number n via the interrupt controller.
 - Saves CS, EIP, EFLAGS, SS, and ESP onto the stack.
 - Jumps to the "Gate" in `IDT[n]`.
 - **A data structure that describes privilege-level switching and long jumps.**
- **RISC-V (M-Mode, Direct Exception Mode)**
 - Checks whether this interrupt should be masked.
 - Jumps: `PC = (mtvec & ~0xF)`
 - Updates: `mcause.Interrupt = 1`

Forcibly "Injected" syscalls

- **Interrupt**

- Saves: `mepc = PC`
- Jumps: `PC = (mtvec & ~0xF)`
- Updates: `mcause.Interrupt = 1`

- **System Call (ecall)**

- Saves: `mepc = PC`
- Jumps: `PC = (mtvec & ~0xF)`
- Updates: `mcause.Ecall = 1`

"No matter what you are doing right now, go execute the system core code!"

Operating System Kernel (Code)

- Can enable and disable interrupts at will.

User Applications

- Sorry, no direct control over interrupts.
 - You can inspect the flags register (`FL_IF`) in `gdb`.
 - **CLI - Clear Interrupt Flag**
 - `#GP(0)` occurs if **CPL** is greater than **IOPL** and less than 3.
 - Try using: `asm volatile ("cli");`
- Regardless of what code you write, it will always be interrupted.

Assume an Interrupt Occurs

What Should the Operating System Code Do?

- `mov (kernel_rsp), %rsp`
 - This can be fatal.
 - The process (state machine) state will be lost forever.

First: Save the State Machine (Registers)

- Preserve control over memory and data.
- Save register states to ****physical memory**** for later restoration.

Then: Execute the Operating System Code

- C code can freely use registers.
- The OS code selects a state machine for return.
- Restore register states from ****physical memory****.
- Execute `sysret (iret)`.

This is the most elegant pieces of code in operating systems.

Operating System Implementation Tricks

- Set up a "current context".
- Save and restore register states.
 - **AbstractMachine** already helps you obtain registers.

```
Context *on_interrupt(Event ev, Context *ctx) {  
    // Save context.  
    current->context = *ctx;  
  
    // Thread schedule.  
    current = current->next;  
  
    // Restore current thread's context.  
    return &current->context;  
}
```


System call instructions are a special type of “long jump”

- The jump target is pre-configured by the OS and **cannot be controlled by applications.**

Processor interrupts also trigger long jumps to the OS kernel

- The OS kernel **preserves** the process state machine:
 - **Memory pages remain unchanged.**
 - Carefully designed code ensures all registers are safely stored in memory.

At this moment, the system is in a state where:

- **All programs are suspended,** and only OS code is executing.
- The OS selectively schedules the next register context onto the CPU to achieve **context switching.**