

Lecture 19: File System API

(Namespace, Directory, FAT, ext2)

Xin Liu

Florida State University

xl24j@fsu.edu

COP 4610 Operating Systems

<https://xinliulab.github.io/cop4610.html>

November 13, 2024

- Abstraction of I/O devices
 - Device layer: I/O devices (registers and protocols)
 - Driver layer: read / write / ioctl
 - Physical layer: how to use magnetic / pit / electrical to store 1 bit of data
 - Block device layer: block read/write

Question to Answer in This Lecture

- **Q:** How can applications share storage devices?

Main Content of This Lecture

- File system APIs
 - Namespace
 - Directory
- Classic File Systems
 - File Allocation Table (FAT)
 - Second Extended Filesystem (ext2)

Why do we need file system?

Sharing Devices Among Applications

- Multiple processes reading / writing concurrently can lead to conflicts.
 - Race condition
 - A single program bug could compromise the entire operating system
- Should all applications share a drive?

Design Goals of File System:

- 1 Provide a reasonable API for multiple applications to share data
- 2 Offer some level of isolation to ensure that malicious or erroneous programs cannot cause widespread harm

Virtualization of "Storage Device (Byte Sequence)"

- Drive (I/O device) = a readable/writable byte sequence
- **Virtual Drive** (file) = a dynamically readable/writable byte sequence
 - **Namespace Management**
 - Naming, indexing, and traversal of virtual drives
 - **Data Management**
 - `std::vector<char>` (random access/write/resize)

Virtual Drvie: Namespace Management

Virtual Drive: Namespace Management

Organizing Information: Structure virtual drives (files) into a hierarchical system.

Key Points:

- Organize virtual drives (files) into a hierarchical structure for easy access.
- Enable efficient retrieval of data by maintaining logical order.
- Example: Similar to a library categorization system, files can be arranged based on names or categories for quick access.



A library categorization system example, which helps in quick location and retrieval of books, analogous to file system organization.

Directory Tree

- Store logically related data in nearby directories
 - Using Locality of Information

File System "Root"

- **Windows:** Each device (driver) is a separate tree
 - New drive letters are assigned for new devices
 - Simple, direct, convenient, but can be cumbersome (e.g., `game.iso`)
- **UNIX/Linux**
 - Only one root, /
 - What about the second device?



An early computer setup demonstrating the concept of distinct device roots in legacy systems.

Mounting in Directory Trees

UNIX: Allows any directory to be **mounted** as a representation of a device's directory tree.

- Highly flexible design:
 - Devices can be mounted anywhere in the desired location.
 - "Mount points" during Linux installation:
 - `/`, `/home`, `/var` can each be separate drive devices.

Mount System Call

```
int mount ( const char *source, const char *target, const char *filesystemtype,  
            unsigned long mountflags, const void *data );
```

- Example: `mount /dev/sdb /mnt`
- Linux `mount` tool can automatically detect the filesystem

Mounting a File

Mounting a file introduces an interesting loop:

- File = Virtual drive on a hard drive
- Mounting a file = Mounting a virtual drive on a virtual drive

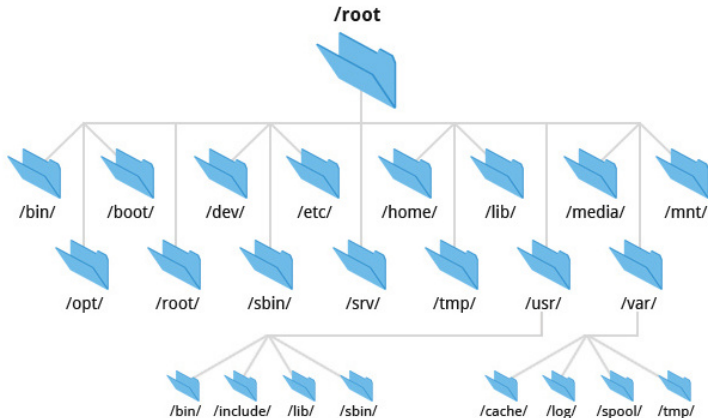
Linux handling:

- Create a **loopback** device
 - Device driver translates device read/write operations to file read/write operations

Filesystem Hierarchy Standard (FHS)

FHS enables *software* and *users* to predict the location of installed files and directories.

- Example: macOS has a UNIX kernel (BSD) but does not follow the Linux FHS.



Directory API (System Calls)

Directory Management: Create/Delete/Traverse

This is straightforward:

- `mkdir`
 - Creates a directory
 - Allows setting access permissions
- `rmdir`
 - Deletes an empty directory
 - No system call for "recursive delete"
 - (If achievable at the application level, it is not implemented at the OS level)
 - `rm -rf` traverses directories, deleting each item (try `strace`)
- `getdents`
 - Returns `count` number of directory entries (used by `ls`, `find`, `tree`)
 - Dot-prefixed entries are returned by the system call, but `ls` does not display them by default

More User-Friendly Directory Access

Appropriate API + Programming Language

- **Globbing**
- This is a user-friendly approach
 - C++ filesystem API is quite difficult to use

Hard Links

Requirements: The system may have multiple versions of the same library.

- Examples: `libc-2.27.so`, `libc-2.26.so`, ...
- Also requires a "current version of libc"
 - Programs need to link to `libc.so.6` to avoid duplicating the file.

Hard Link: Allows a file to be referenced by multiple directory entries.

- Directories only store pointers to the file data.
- **Limitations:**
 - Cannot link directories
 - Cannot link across file systems

Most UNIX file systems use hard links for files (check with `ls -li`).

- System call to delete a link is `unlink` (reference count).

Symbolic Link: Stores a "jump pointer" in a file.

- Symbolic links are also files.
 - When referencing this file, it points to another file.
 - Stores the absolute/relative path of another file as text in the file.
 - Can link across file systems, can link directories, etc.
- Similar to a "shortcut."
 - It doesn't matter if the linked target currently exists.
 - Examples:
 - `~/usb ⇒ /media/xinliu-usb`
 - `~/Desktop ⇒ /mnt/c/Users/xinliu/Desktop (WSL)`

`ln -s` to create symbolic links.

- `symlink` system call.

Process "Current Directory"

Working/Current Directory

- `pwd` command or `$PWD` environment variable can be used to check.
- `chdir` system call for modification.
 - Corresponds to `cd` in the shell.
 - Note that `cd` is a shell built-in command.
 - It does not exist in `/bin/cd`.

Question: Do threads share a working directory, or does each have its own?

File API (System Calls)

Review: Files and File Descriptors

Files: Virtual Drives

- A drive is a "sequence of bytes."
- Supports read/write operations.

File Descriptors: Pointers for Process Access to Files (Operating System Objects)

- Obtained through `open` or `pipe`.
- Released through `close`.
- Duplicated through `dup/dup2`.
- Inherited during `fork`.

File Access Offset (Seek Pointer)

File read/write operations come with a "seek pointer," so it's unnecessary to specify the read/write location every time.

- This feature makes it convenient for programmers to access files sequentially.

Example:

- `read(fd, buf, 512);` - Reads the first 512 bytes.
- `read(fd, buf, 512);` - Reads the next 512 bytes.
- `lseek(fd, -1, SEEK_END);` - Moves to the last byte.
 - *so far, so good*

Offset Management: Not So Simple

File descriptors are inherited by child processes during `fork`.

Should parent and child processes share an offset, or should each have its own?

- This choice determines where the offset is stored.

Consider application scenarios:

- When parent and child processes write to a file simultaneously
 - Each has its own offset → parent and child need to coordinate offset updates
 - (Race condition)
 - Shared offset → OS manages the offset
 - Although shared, the OS ensures the atomicity of `write` operations ✓

Offset Management: Behavior

Every API in the operating system may interact with other APIs

- 1 During `open`, a unique offset is obtained.
- 2 During `dup`, two file descriptors share the offset.
- 3 During `fork`, the parent and child processes share the offset.
- 4 During `execve`, the file descriptor remains unchanged.
- 5 For files opened with `O_APPEND` mode, the offset is always at the end (regardless of `fork`).
 - Modification of the file offset and the write operation are performed as a single atomic step.

This is also one reason why `fork` is often criticized.

- (At the time) a good design may become a burden in the evolution of the system.
- Today's `fork` might be considered "overloaded"; *A fork() in the road.*

File Allocation Table (FAT)

What is File System Implementation?

- Implement all file system APIs on a block device (I/O device)
 - `bread(int id, char *buf);`
 - `bwrite(int id, const char *buf);`
 - Assumes all operations complete in synchronized queue
 - (Can be implemented with queues at block I/O layer)

Directory/File API

- `mkdir`, `rmdir`, `link`, `unlink`
- `open`, `read`, `write`, `stat`

Back to Data Structures Class...

- A file system is essentially a data structure (Abstract Data Type; ADT)
 - Just with different assumptions than those in data structures class

Assumptions in Data Structures Class:

- Von Neumann machine
- Random Access Memory (RAM)
 - Word Addressing (e.g., 32/64-bit load/store)
 - The cost of each instruction is $O(1)$
 - Memory hierarchy challenges this assumption (cache-unfriendly code may encounter performance issues)

Assumptions in File Systems:

- Block-based (e.g., 4KB) access; building a RAM model on disk is entirely unrealistic

Device Abstraction Provided by Block Device

Implementation of Data Structures (cont'd)

- Virtualizing the disk with `balloc/bfree`
 - `File = vector<char>`
 - Maintains using linked lists, indexes, or any data structure
 - Supports arbitrary position modifications and resizing
- Implementing directories based on files
 - `Directory file`
 - Interprets `vector<char>` as `vector<dir_entry>`
 - Stores continuous bytes for each directory entry

Implementation of a Simple File System

We can implement a simple file system by treating files as sequences of blocks and using data structures to manage them.

- **File Representation**

- Each file is represented by an **inode** (index node)
- The inode contains metadata and pointers to data blocks

- **Inode Structure**


- File type, permissions, owner, timestamps, etc.
- Pointers to data blocks (direct, indirect, double indirect)

- **Data Blocks**

- Fixed-size blocks storing the actual file content
- Managed using block allocation algorithms

Back to 1980: The 5.25" Floppy Disk

- **5.25" Floppy Disk:** Single-sided, 180 KiB capacity
- **Storage Specifications:**
 - 360 sectors, each with 512 bytes (sectors)
- **Question:**
 - What kind of data structure would be suitable to implement a file system on such a device?



Files in the FAT File System

- **Characteristics:**

- Relatively small file system
- Tree-like directory structure
- Primarily consists of small files (within a few blocks)

- **File Implementation:**

- Linked list of `struct block *`
- Complex high-level data structures are inefficient for this purpose


Using Linked Storage Data: Two Designs

- 1 Place the pointer after each data block
 - **Advantage:** Simple implementation, no need for additional storage space.
 - **Disadvantage:** Data size is not necessarily 2^k ; pure `lseek` requires reading entire block data.
- 2 Centralize pointers in a specific area of the file system
 - **Advantage:** Better locality; faster `lseek`.
 - **Disadvantage:** Centralized pointer data corruption could lead to data loss.

Question: Which design's drawbacks are fatal and difficult to resolve?

Centralized Storage of All Pointers

- Centralized pointers are prone to damage? Store n copies to mitigate!
- Example: FAT-12/16/32 (FAT entry represents the size of the "next pointer")



FAT: Linked Storage Files

- Structure of FAT's "next" array:
 - 0: free; 2 ... MAX: allocated;
 - 0xFFFFFFFF7: bad cluster; 0xFFFFFFFF8 - 0xFFFFFFFEE, -1: end-of-file

FAT32	Comments
0x00000000	Cluster is free.
0x00000002 to MAX (MAX + 1) to 0xFFFFF6	Cluster is allocated. Value of the entry is the cluster number of the next cluster following this cluster.
0xFFFFF7	Reserved and must not be used.
0xFFFFF8 to 0xFFFFFEE	Indicates a bad (defective) cluster.
0xFFFFFFFF	Reserved and should not be used.
	Cluster is allocated and is the final cluster for the file (indicates <i>end-of-file</i>).

Table: FAT Entry Values and Their Meanings

Directory Tree Implementation: Directory Files

Using regular files to store the "directory" data structure

- **FAT:** Directory is a collection of fixed-length 32-byte directory entries.
- The operating system parses and treats directory entries marked as "directory" as actual directories.
 - A sequence of directory entries can store long filenames.
- **Thought Exercise:** Why not store metadata (size, filename, etc.) at the head of `vector<struct block *> file`?

FAT: Performance and Reliability

Performance

- + Small files are ideal
- - However, random access for large files is inefficient
 - A 4 GB file jumping to the end (4 KB clusters) requires 2^{20} chain 'next' operations
 - Caching can partially alleviate this issue
- In the FAT era, sequential access performance on disks was better
 - Long-term disk usage leads to fragmentation
 - `malloc` also causes fragmentation, but the performance impact is less significant

Reliability

- Maintain multiple copies of FAT to prevent data loss
 - Unexpected synchronous write-offs
 - Damaged clusters are marked in the FAT

ext2 and UNIX File System

Centralized storage of file/directory metadata as objects

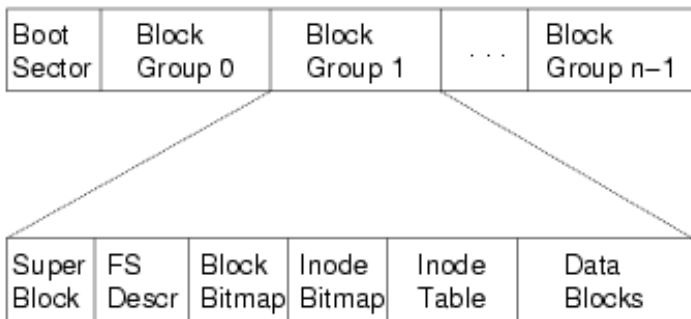
- Enhances locality (easier for caching)
- Supports linked files

Distinguishing fast/slow paths for different file sizes

- For small files, arrays should be used
 - Skips linked list traversal
- For large files, trees should be used (e.g., B-Tree, Radix-Tree)
 - Enables fast random access

ext2: Disk Image Format

Dividing the disk into groups



Superblock: Filesystem metadata

- Number of files (inodes)
- Block group information
 - `ext2.h` contains everything you need to know

ext2 Directory Files

Similar to FAT: Establishes a directory structure on files

- Note that inodes are stored in a unified way
 - Directory files store a key-value mapping of file names to inode numbers

ext2: Performance and Reliability

For large files, random read/write performance is significantly improved ($O(1)$):

- Supports linking (reduces space waste to some extent)
- Inodes are stored continuously on disk, which facilitates caching/prefetching
- Fragmentation remains an issue

However, reliability is still a major concern:

- Damage to the data block storing the inode can be very serious

This Lecture's Key Question

- **Q:** How to design a file system that allows applications to share storage devices?

Takeaway Messages

- Two Main Components of File System
 - Virtual Drive (File)
 - Functions: `mmap`, `read`, `write`, `lseek`, `ftruncate`, ...
 - Virtual Drive Naming and Management (Directory Tree and Links)
 - Functions: `mount`, `chdir`, `mkdir`, `rmdir`, `link`, `unlink`, `symlink`, `open`, ...